

# Midterm-Rom-P

October 8, 2020

## 1 CP - Midterm - 2020

### 1.1 Instruction

- Modify this file to be Midterm-, e.g., Midterm-Chaklam-S.ipynb
- This exam accounts for 25% of the overall course assessment.
- This exam is open-booked; open-internet.
- You ARE NOT allowed to use sklearn or any libraries, unless stated.
- The completed exams shall be submitted at the Google Classroom
- All code should be **complemented with comments**, unless it's really obvious. **I and Joe reserve the privilege to give you zero for any part of the question where the benefit of doubt is not justified**

### 1.2 Examination Rules:

- For **offline** students, you may leave the room temporarily with the approval and supervision of the proctors. No extra time will be added to the exam in such cases.
- For **online** students, you are required to turn on your webcam during the entire period of the exam time
- Students will be allowed to leave at the **earliest 45 minutes** after the exam has started
- **All work should belong to you.** A student should NOT engage in the following activities which proctors reserve the right to interpret any of such act as academic dishonesty without questioning:
  - Chatting with any human beings physically or via online methods
  - Plagiarism of any sort, i.e., copying from internet sources or friends. **Both copee and copier shall be given a minimum penalty of zero mark for that particular question or the whole exam.**
- No make-up exams are allowed. Special considerations may be given upon a valid reason on unpredictable events such as accidents or serious sickness.

### 1.3 Question 1 (21 pts)

#### 1). The rabbit: (5pts)

Once upon a time, there is a father rabbit lives in a far away jungle. Everyday, the father rabbit has to go out and find some carrots for his family. In his family there are mother rabbit, grampa rabbit, sister rabbit, and his son. In total there are 5 rabbits to feed. In one day, the adult rabbits

(himself, mother rabbit and sister rabbit) will eat 3 carrots while the elderly eat 2 carrots and baby rabbit eat 1 carrot.

Unfortunately, the carrots are not easy to find. The father rabbit has to travel into the scary jungle and find some carrot then bring them back to the family before the sunset at 6PM.

- Every 1 km, the rabbit will find 3 carrots.
- The rabbit will use 1 hour to travel 1 km.

In summary, in order to find the least number of carrot for each day, the rabbit will have to use  $(3 + 3 + 3 + 2 + 1)/3 = 4$  hours. This mean that he has to leave the house at the latest 10AM (4 hours for go out and another 4 for comming back).

This daily work has to be done exactly on time, leaving to late will cause whether his life or his family life. Would you like to help the rabbit?

```
[1]: # print 'yes' to help the rabbit or 'no' to refuse the challenge. (if yes -> 1,
      ↪pt)

      #Father Rabbit 3 carrots
      #Mother Rabbit 3 carrots
      #Grandpa Rabbit 2 carrots
      #Sister Rabbit 3 carrots
      #Son Rabbit 1 carrots

      #Leaves house at the latest 10AM
      #yes
```

Good to hear that young programmer!!

What I have in mind is to build a clock that when the rabbit puts the number of (adult, elderly, young) rabbit, it will calculate how many hours is required for a travel. Of cause we have to make it as a function because the number of each rabbit type will change over the time.

- Write a function carrot that takes three integers as an input in follow this format (adult, elderly, young) (1pt)
- The function will calculate number of hour required for travelling a day. (1pt)
- The function will also calcualte the time to leave. Think of it as an alarm clock for leaving the house (1pt)
- The function will return a tuple (#hours, #time) (1pt)

```
[2]: # Your code here

def carrot(adult, elderly, young):
    hours = (adult*3+elderly*2+young)/3
    leave = 18-2*hours
    #     if(hours < 12):
    #         print(leave)
    #         hour_str = str(int(leave))+":"+str(int(60*(leave%1)))+ "AM"
    #     else:
    #         hour_str = str(int(leave-12))+":"+str(int(60*(leave%1)))+ "PM"
```

```

    return (hours,leave)

carrot(3,1,1)

```

[2]: (4.0, 10.0)

2). **Print the shape:** (16pts)

- Write a function square that takes integer as an input. (1pt)
- The function will return a string of \* in the shape of a square with both width and height equal to the input integer. (2pts)

```

[3]: '''
Level: 3
***
***
***

Level: 5
*****
*****
*****
*****
*****
'''

```

[3]: '\nLevel: 3\n\*\*\*\n\*\*\*\n\*\*\*\n\nLevel: 5\n\*\*\*\*\*\n\*\*\*\*\*\n\*\*\*\*\*\n\*\*\*\*\*\n\*\*\*\*\*\n'

```

[4]: # Your code here
def square(n):
    #Create two loops to print "*"
    for j in range(n):
        for i in range(n):
            print("*",end="")
            print()

square(5)

```

```

*****
*****
*****
*****
*****

```

- Write a function triangle that takes integer as an input. (1pt)
- The function will return a string of \* in the shape of a triangle with level equal to the input integer. (2pts)

```
[5]: '''
Level: 3
    *
   **
  ***

Level: 5
    *
   **
  ***
 ****
*****
'''
```

```
[5]: '\nLevel: 3\n    *\n   **\n  ***\n\nLevel: 5\n    *\n   **\n  ***\n ****\n*****\n'
```

```
[6]: # Your code here

def triangle(n):
    #Create a loop for each row, each row will have its corresponding number of
    ↳spaces and number of "*"
    for i in range(n):
        spaces = " "(n-i-1)
        ast = "*"*(i+1)
        print(spaces,end="")
        print(ast)

triangle(3)
```

```
*
**
***
```

- Write a function pyramid that takes integer as an input. (1pt)
- The function will return the string of \* in the shape of a pyramid with level equal to the input interger. (2pts)

```
[7]: '''
Level: 3
    *
   **
  ***
 ****

Level: 5
    *
   **
  ***
 ****
*****
'''
```

```

*****
*****
'''

```

```

[7]: ' \nLevel: 3 \n * \n ***\n*****\n\nLevel: 5\n      *\n      ***\n      *****\n      *****\n      *****\n      *****\n'

```

```

[8]: # Your code here

```

```

def pyramid(n):
    #similar to the previous question, each row will have their respective
    ↳ amount of spaces and asterisks
    for i in range(n):
        spaces = " "(n-i-1)
        ast = "*"*(i*2+1)
        print(spaces,end=" ")
        print(ast,end=" ")
        print(spaces)

pyramid(5)

```

```

*
***
*****
*****
*****

```

Now, let's combine the three algorithms into one single class. - Create a class named MyShape that can do the followings - Take two arguments during the class construction. The first one is an integer and the second one is a string. The names are level and shape (1pt) - Check the input arguments whether the integer is in the range of [1,10] and string is in the set of {'squ','tri','pyr'}. Raise a ValueError. (2pts) - Both attributes should be able to change via a set method **only**. set[attrName] (1pt) - Of course, the set method should check the out of range too. (1pt) - To check the current setting, write a get method. get[attrName] (1pt) - Print the shape with method show. It should return the string of the current shape with the correct level (1pt)

```

[9]: '''
Example 1

>>> ms = MyShape(2,'tri')
>>> ms.show()

*
**
>>> ms.setLevel(3)
>>> ms.setShape('squ')
>>> ms.show()

***
***
'''

```

```

***
>>> ms.setShape('a')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: .....
>>>
'''

```

```

[9]: '\nExample 1\n\n>>> ms = MyShape(2,\n'tri\n')\n>>> ms.show()\n * \n**\n>>>
ms.setLevel(3)\n>>> ms.setShape('\nsqu\n')\n>>> ms.show()\n***\n***\n***\n>>>
ms.setShape('\na\n')\nTraceback (most recent call last):\n  File "<stdin>", line
1, in <module>\nValueError: ... \n>>>\n'

```

```

[10]: # Your code here
class MyShape():
    def __init__(self,level,shape):
        self.l = level
        self.s = shape
        self.checkLevel(self.l)
        self.checkShape(self.s)

    def checkLevel(self,level):
        if(level > 10 or level < 1):
            raise ValueError("Invalid Level")

    def checkShape(self,shape):
        if(shape not in ["squ","tri","pyr"]):
            raise ValueError("Invalid Shape")

    def setLevel(self,level):
        self.checkLevel(level)
        self.l = level

    def setShape(self,shape):
        self.checkShape(shape)
        self.s = shape

    def getLevel(self):
        return self.l

    def getShape(self):
        return self.s

    def pyramid(self,n):
        for i in range(n):
            spaces = " "*(n-i-1)
            ast = "*"*(i*2+1)

```

```

        print(spaces,end="")
        print(ast,end="")
        print(spaces)

    def triangle(self,n):
        for i in range(n):
            spaces = " "(n-i-1)
            ast = "*"*(i+1)
            print(spaces,end="")
            print(ast)

    def square(self,n):
        for j in range(n):
            for i in range(n):
                print("*",end="")
            print()

    def show(self):
        if(self.s == 'squ'):
            self.square(self.l)
        elif(self.s == 'tri'):
            self.triangle(self.l)
        elif(self.s == 'pyr'):
            self.pyramid(self.l)
        else:
            print("HI")

ms = MyShape(3,'pyr')
print(ms.getShape())
print(ms.getLevel())
ms.show()
ms.setShape('squ')
ms.setLevel(10)
print(ms.getShape())
print(ms.getLevel())
ms.show()

```

```

pyr
3
  *
***
*****
squ
10
*****
*****

```

```
*****
*****
*****
*****
*****
*****
*****
*****
```

## 1.4 Question 2 (10 pts)

### 2). ML Skill

$$y = ax + b$$

The above equation is your favorite linear equation where a,b are the constant value indicate the slope and the offset of the line in the graph.

We all know given and two points.

$$(x_1, y_1)(x_2, y_2)$$

you can find a,b very easy using Geometry

$$a = \frac{y_2 - y_1}{x_2 - x_1}$$

$$b = y_i - ax_i$$

Since we have learnt that using LinearRegression can find the value of the a,b too.

Now, do the followings.

- Write a function drawLine that takes two tuples as inputs.
- Calculate a,b using Geometry.
- Draw the first graph with scatter on the given two points and a line.
- Calculate a,b using LinearRegression with Batch Gradient Descent.
  - Generate 1000 sample data along the line.
  - Regress on the data using LinearRegression-BatchGradientDescent
- Draw the second graph with scatter on the given two points and a line.
- Does both method yeild the same outcome? Which method runs faster? (use timeit)
- What will happen if the data is normalize first? (draw another graph and timeit)
- What will happen if the data is standardize first? (draw another graph and timeit)

```
[11]: # Your code here
import matplotlib.pyplot as plt
import numpy as np

class LinearRegressionModel:
    #1. hypothesis function
    def h(self, X, theta):
        hypothesis = X@theta
        return hypothesis
```



#2. cost function

```
def cost(self, X, y, theta, average = False):
    #expects X to be a design matrix, y to be a column vector and theta to
    → be a column vector
    if(average == False):
        J = 1/2*(self.h(X,theta)-y).T@(self.h(X,theta)-y)
    else:
        J = 1/(2*X.shape[0])*(self.h(X,theta)-y).T@(self.h(X,theta)-y)
    return J
```

#3. gradient function

```
def gradient(self, X, y, theta, average = False):
    if(average == False):
        dJ = X.T@(self.h(X,theta)-y)
    else:
        dJ = X.T@(self.h(X,theta)-y)/(X.shape[0])
    return dJ
```

#4. batch gradient descent

```
def batch_gd(self, X, y, initial_theta, max_iteration, alpha, tolerance =
    → 0, average = False):
    cost = []
    theta = initial_theta
    iteration = 0
    cost.append(self.cost(X,y,theta,average))
    for n in range(max_iteration):
        gradient = self.gradient(X,y,theta,average)
        theta = theta - alpha*gradient
        cost.append(self.cost(X,y,theta,average))
        iteration += 1
        if(self.mean_squared_error(X,y,theta) < tolerance):
            return theta,cost,iteration
    cost = np.array(cost)
    return theta,cost,iteration
```

#5. normal equation

```
def normal_equation(self, X, y):
    theta = np.linalg.inv(X.T@X)@X.T@y
    return theta
```

#5. predict

```
def predict(self,X,theta):
    prediction = self.h(X,theta)
    return prediction
```

#6. score/error calculation

```

def mean_squared_error(self,X,y,theta):
    mse = self.cost(X,y,theta,average = True)*2
    return mse

#7. plotting cost
def plot_cost(self,cost, iteration_no):
    iteration_series = np.arange(0,iteration_no+1)
    ax = plt.axes()
    ax.plot(iteration_series, cost)

def drawLine(one,two):
    # drawLine function as stated in the problem
    x,y = zip(one,two)
    print("timeit for using the equations")
    #using timeit, we can see how long it takes to perform these operations
    a = (y[1]-y[0])/(x[1]-x[0])
    b = y[0]-a*x[0]
    %timeit (y[1]-y[0])/(x[1]-x[0]), y[0]-a*x[0]

    #Plot for equation
    fig,ax = plt.subplots(2,2,figsize = (20,20))
    ax[0,0].scatter(x,y)
    ax[0,0].plot(x,y, 'r')

    x_sample = np.linspace(np.min(x),np.max(x),1000)
    y_sample = x_sample*a+b

    LR = LinearRegressionModel()
    iterations = 1000
    alpha = 0.001
    initial_theta = np.zeros(2)
    x_sample_inserted = np.insert(x_sample[:,np.newaxis],0,1,axis=1)

    #using batch gradient descent on data which hasn't been standardized or
    ↪normalized
    print("timeit for batch gradient descent")
    %timeit LR.
    ↪batch_gd(x_sample_inserted,y_sample,initial_theta,iterations,alpha)
        theta,cost,iteration = LR.
    ↪batch_gd(x_sample_inserted,y_sample,initial_theta,iterations,alpha,tolerance=1e-7)
        y_pred = LR.predict(x_sample_inserted,theta)
        #LR.plot_cost(cost,iteration)
        print("iteration:",iteration)

    MSE = LR.mean_squared_error(x_sample_inserted,y_sample,theta)
    print("MSE =",MSE)

```

```

ax[0,1].scatter(x,y)
ax[0,1].plot(x,y,'r')
ax[0,1].scatter(x_sample,y_pred)

mini = np.min(x_sample)
maxi = np.max(x_sample)
x_norm = (x_sample-mini)/(maxi-mini)

x_norm_inserted = np.insert(x_norm[:,np.newaxis],0,1,axis=1)

iterations2 = 1000
alpha2 = 0.001
initial_theta2 = np.zeros(2)

print("timeit for normalized data")
%timeit LR.
↪batch_gd(x_norm_inserted,y_sample,initial_theta2,iterations2,alpha2)
    theta2,cost2,iteration2 = LR.
↪batch_gd(x_norm_inserted,y_sample,initial_theta2,iterations2,alpha2,tolerance=1e-7)
    y_pred2 = LR.predict(x_norm_inserted,theta2)
    print("iteration2",iteration2)
    MSE2 = LR.mean_squared_error(x_norm_inserted,y_sample,theta2)
    print("MSE =",MSE2)

ax[1,0].scatter(x,y)
ax[1,0].plot(x,y,'r')
ax[1,0].scatter(x_norm*(maxi-mini)+mini,y_pred2)

mean = np.mean(x_sample)
std = np.std(x_sample)
x_stan = (x_sample-mean)/std

x_stan_inserted = np.insert(x_stan[:,np.newaxis],0,1,axis=1)

iterations3= 1000
alpha3 = 0.001
initial_theta3 = np.zeros(2)

print("timeit for standardized data")
%timeit LR.
↪batch_gd(x_stan_inserted,y_sample,initial_theta3,iterations3,alpha3,average=True)
    theta3,cost3,iteration3 = LR.
↪batch_gd(x_stan_inserted,y_sample,initial_theta3,iterations3,alpha3,tolerance=1e-7)
    y_pred3 = LR.predict(x_stan_inserted,theta3)
    print("iteration3",iteration3)
    MSE3 = LR.mean_squared_error(x_stan_inserted,y_sample,theta3)
    print("MSE =",MSE3)

```

```

ax[1,1].scatter(x,y)
ax[1,1].plot(x,y,'r')
ax[1,1].scatter(x_stan*std+mean,y_pred3)

print("slope =",a,"intercept =",b)


drawLine((1,2),(-0,0))

```

```

timeit for using the equations
212 ns ± 13.3 ns per loop (mean ± std. dev. of 7 runs, 1000000 loops each)
timeit for batch gradient descent
28.2 ms ± 1.26 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)
iteration: 107
MSE = 9.514767563793411e-08
timeit for normalized data
26 ms ± 260 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)
iteration2 107
MSE = 9.514767563793411e-08
timeit for standardized data
29 ms ± 1.38 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)
iteration3 1
MSE = 6.848366140997965e-32
slope = 2.0 intercept = 0.0

```



### 1.5 Question 3 (69 pts)

1). **Exploratory Data Analysis:** - Load the data "howlongwelive.csv" to pandas and print the first 5 and last 5 rows of data (1 or 0pt)

- Print the shape, feature names, and summary (describe) of the data (1 or 0pt)
- Check whether there is missing data. (1 or 0pt)
- Fix all missing data using means or mode (1 or 0pt)
- Since Hepatitis B has a lot of nans, and highly correlate with Diptheria, simply drop column Hepatitis. Also drop column Population since there are way too many nans (1 or 0pt)
- If there are any features which are string and you want to use them as features, we need to convert them to int or float. For now, convert Status to 0 or 1 (1 or 0pt)

- Rename column `thinness_1-19_years` to `thinness_10-19_years` (1 or 0pt)
  - Perform a groupby country and plot their life expectancy. Which country has the lowest/highest life expectancy? (1 or 0pt)
  - Plot average life expectancy of developed country vs. developing country. (1 or 0pt)
  - Perform a t-test of life expectancy between developed and developing countries. Is the result significant? (1 or 0pt)
  - Perform a pairplot to see which features are likely to have strong predictive power for life expectancy. Identify the most 3 important features. (1 or 0pt)
  - Perform a histogram of life expectancy. Is it normal? (1 or 0pt)
- 2). **Regression** - Prepare your X and y into Numpy array (you have to map from Pandas to numpy). For X, prepare two versions of them. For first `X_selected`, you have to choose the most 3 important features from above, and for second `X_all`, simply use all features (you may want to omit Country since they are categorical). Set y to life expectancy. (1 or 0pt)
- Perform standardization using Numpy way (NOT sklearn way). (1 or 0pt)
  - Perform train-test split by using Numpy way (NOT sklearn way). Use test size of 0.3. (1 or 0pt)
  - Perform assertion whether your splitting is correct accordingly (1 or 0pt)
  - Write a class `Regression(X, y, grad_method, max_iter, alpha, tol, decay, decay_iter, decay_rate, stop_decay_counter, verbose, lam, poly, poly_deg)` that can perform the followings:
    - Mini-batch, Stochastic, and Batch Gradient Descent (each 2pts)
    - Polynomial of degree k (2 or 0pt)
    - Decay learning rate (1 or 0pt)
      - \* Decay learning rate is a learning rate that becomes smaller after certain iteration. For example, after 5 iterations, the learning rate will reduce to 95% of the current learning rate.
      - \* To implement it, simply multiply current learning rate with some constant `decay_rate`. For now, set it to 0.9
    - Regularization with ridge (2 or 0pt)
    - Must have at least four methods for `fit()` (i.e., for finding weights) `predict()` (i.e., for predicting `X_test` data), `score()` (i.e., for returning  $r^2$  score), and `mse()` (return mse) (each 1pt)
    - Accepts X, y, `grad_method` (default set to "batch"), `alpha` (learning rate), `max_iter`, `tol`, `decay` (whether to use decay learning rate; default set to False), `decay_iter` (after how many iterations will the decay apply), `stop_decay_counter` (this is the maximum number of times that decay the learning rate), `verbose` (default is set to False, whether model will display the Cost for each iteration), `lam` (this is the ridge regularization parameter), `poly` (default is set to False), and `poly_deg` (default is set to 2) (each 1/13pt)
  - Create the following 3 models **from your class** (For any unspecified parameters, feel free to use any :D)

1. For the first model, transform your feature using polynomial degree 3, then perform linear regression with batch gradient descent with early stopping of tol 1e-3 (1 or 0pt)
  2. For the second model, perform linear regression with mini-batch gradient descent with early stopping of tol 1e-3 (1 or 0pt)
  3. For the third model, perform ridge regression with stochastic gradient descent with early stopping of tol 1e-3 and decay set to True and lam to 1e-4 (1 or 0pt)
- Create Lasso model from Sklearn with default parameters (1 or 0pt)
  - For these four models, using two different versions of X, perform a cross validation of 10 folds, comparing the four models \* two versions of X. Here you should implement cross validation. Report which one is the best candidate model (3pts for implement from scratch or 1pt for using sklearn)
    - Recall that in a 10 folds cross validation, you split your data into 10 even pieces. Then you run 10 iterations, where in each iteration, you pick 1 of this piece as the validation set, and the rest as training set. Once you reach the 10th iteration, you would have already exhaust all the 10 pieces as validation set.
  - Using the best model, fit again with the training data. Plot the weights using bar charts along the feature names. Before you actually plot the weights, we need to multiply these weights by their feature standard deviation, so to reduce these weights to same unit of measure. Interpret these weights and what they imply. (For those who are curious why we need to multiply with std, you may read this > [https://scikit-learn.org/stable/auto\\_examples/inspection/plot\\_linear\\_model\\_coefficient\\_interpretation.html#interpretation-coefficients-scale-matters](https://scikit-learn.org/stable/auto_examples/inspection/plot_linear_model_coefficient_interpretation.html#interpretation-coefficients-scale-matters) (2 or 0pt)
  - Perform predictions on testing data. Print adjusted  $r^2$  and mse. (1 or 0pt)
  - Plot the predicted values against actual values (1 or 0pt)

### 3). Classification

- Change your y to discrete value. Here split y into three class, {0, 1, 2}, where 0 belongs to low life expectancy group, and 2 for the high life expectancy group. (1 or 0pt)
- Write a class for multinomial logistic regression with stochastic gradient descent. Must have at least six methods for fit() (i.e., for finding weights) predict() (i.e., for predicting X\_test data), accuracy() (i.e., for returning accuracy score), recall(), precision(), and f1() (each 1pt)
- Using the best X\_train of the two suggested by the cross validation step, fit the data with your class. (1 or 0pt)
- Perform predictions on testing data. Print accuracy, recall, precision, and f1\_score from your class. (1 or 0pt)
- Plot the decision boundary with the X\_test data. To plot this, you may want to choose only 2 features. (1 or 0pt)

### 4). Final verdict

- Attempt to do whatever ways - including sklearn or scratch - or change your features, or do feature engineering such that your mse is lowest possible. (0 to 5pts - following class normal distributions)

1). **Exploratory Data Analysis:** - Load the data "howlongwelive.csv" to pandas and print the first 5 and last 5 rows of data (1 or 0pt)

- Print the shape, feature names, and summary (describe) of the data (1 or 0pt)
- Check whether there is missing data. (1 or 0pt)
- Fix all missing data using means or mode (1 or 0pt)
- Since Hepatitis B has a lot of nans, and highly correlate with Diphtheria, simply drop column Hepatitis. Also drop column Population since there are way too many nans (1 or 0pt)
- If there are any features which are string and you want to use them as features, we need to convert them to int or float. For now, convert Status to 0 or 1 (1 or 0pt)
- Rename column thinness\_1-19\_years to thinness\_10-19\_years (1 or 0pt)
- Perform a groupby country and plot their life expectancy. Which country has the lowest/highest life expectancy? (1 or 0pt)
- Plot average life expectancy of developed country vs. developing country. (1 or 0pt)
- Perform a t-test of life expectancy between developed and developing countries. Is the result significant? (1 or 0pt)
- Perform a pairplot to see which features are likely to have strong predictive power for life expectancy. Identify the most 3 important features. (1 or 0pt)
- Perform a histogram of life expectancy. Is it normal? (1 or 0pt)

```
[12]: # Your code here
import pandas as pd

data = pd.read_csv("howlongwelive.csv")
```

```
[13]: data.head()
```

```
[13]:
```

	Country	Year	Status	Life expectancy	Adult Mortality	\
0	Afghanistan	2015	Developing	65.0	263.0	
1	Afghanistan	2014	Developing	59.9	271.0	
2	Afghanistan	2013	Developing	59.9	268.0	
3	Afghanistan	2012	Developing	59.5	272.0	
4	Afghanistan	2011	Developing	59.2	275.0	

	infant deaths	Alcohol	percentage expenditure	Hepatitis B	Measles	...	\
0	62	0.01	71.279624	65.0	1154	...	
1	64	0.01	73.523582	62.0	492	...	
2	66	0.01	73.219243	64.0	430	...	
3	69	0.01	78.184215	67.0	2787	...	
4	71	0.01	7.097109	68.0	3013	...	

	Polio	Total expenditure	Diphtheria	HIV/AIDS	GDP	Population	\
0	6.0	8.16	65.0	0.1	584.259210	33736494.0	



1	58.0	8.18	62.0	0.1	612.696514	327582.0
2	62.0	8.13	64.0	0.1	631.744976	31731688.0
3	67.0	8.52	67.0	0.1	669.959000	3696958.0
4	68.0	7.87	68.0	0.1	63.537231	2978599.0

	thinness 1-19 years	thinness 5-9 years \
0	17.2	17.3
1	17.5	17.5
2	17.7	17.7
3	17.9	18.0
4	18.2	18.2

	Income composition of resources	Schooling
0	0.479	10.1
1	0.476	10.0
2	0.470	9.9
3	0.463	9.8
4	0.454	9.5

[5 rows x 22 columns]

```
[14]: data.tail()
```

```
[14]:
```

	Country	Year	Status	Life expectancy	Adult Mortality \
2933	Zimbabwe	2004	Developing	44.3	723.0
2934	Zimbabwe	2003	Developing	44.5	715.0
2935	Zimbabwe	2002	Developing	44.8	73.0
2936	Zimbabwe	2001	Developing	45.3	686.0
2937	Zimbabwe	2000	Developing	46.0	665.0

	infant deaths	Alcohol	percentage expenditure	Hepatitis B	Measles \
2933	27	4.36	0.0	68.0	31
2934	26	4.06	0.0	7.0	998
2935	25	4.43	0.0	73.0	304
2936	25	1.72	0.0	76.0	529
2937	24	1.68	0.0	79.0	1483

	...	Polio	Total expenditure	Diphtheria	HIV/AIDS	GDP \
2933	...	67.0	7.13	65.0	33.6	454.366654
2934	...	7.0	6.52	68.0	36.7	453.351155
2935	...	73.0	6.53	71.0	39.8	57.348340
2936	...	76.0	6.16	75.0	42.1	548.587312
2937	...	78.0	7.10	78.0	43.5	547.358879

	Population	thinness 1-19 years	thinness 5-9 years \
2933	12777511.0	9.4	9.4
2934	12633897.0	9.8	9.9

2935	125525.0	1.2	1.3
2936	12366165.0	1.6	1.7
2937	12222251.0	11.0	11.2

	Income composition of resources	Schooling
2933	0.407	9.2
2934	0.418	9.5
2935	0.427	10.0
2936	0.427	9.8
2937	0.434	9.8

[5 rows x 22 columns]

```
[15]: print("shape:",data.shape)
      print("features:",data.columns)
      data.describe()
```

```
shape: (2938, 22)
features: Index(['Country', 'Year', 'Status', 'Life expectancy ', 'Adult Mortality',
                'infant deaths', 'Alcohol', 'percentage expenditure', 'Hepatitis B',
                'Measles ', ' BMI ', 'under-five deaths ', 'Polio', 'Total expenditure',
                'Diphtheria ', ' HIV/AIDS', 'GDP', 'Population',
                ' thinness 1-19 years', ' thinness 5-9 years',
                'Income composition of resources', 'Schooling'],
              dtype='object')
```

```
[15]:
```

	Year	Life expectancy	Adult Mortality	infant deaths \
count	2938.000000	2928.000000	2928.000000	2938.000000
mean	2007.518720	69.224932	164.796448	30.303948
std	4.613841	9.523867	124.292079	117.926501
min	2000.000000	36.300000	1.000000	0.000000
25%	2004.000000	63.100000	74.000000	0.000000
50%	2008.000000	72.100000	144.000000	3.000000
75%	2012.000000	75.700000	228.000000	22.000000
max	2015.000000	89.000000	723.000000	1800.000000

	Alcohol	percentage expenditure	Hepatitis B	Measles \
count	2744.000000	2938.000000	2385.000000	2938.000000
mean	4.602861	738.251295	80.940461	2419.592240
std	4.052413	1987.914858	25.070016	11467.272489
min	0.010000	0.000000	1.000000	0.000000
25%	0.877500	4.685343	77.000000	0.000000
50%	3.755000	64.912906	92.000000	17.000000
75%	7.702500	441.534144	97.000000	360.250000
max	17.870000	19479.911610	99.000000	212183.000000

	BMI	under-five deaths	Polio	Total expenditure \
count	2904.000000	2938.000000	2919.000000	2712.000000
mean	38.321247	42.035739	82.550188	5.93819
std	20.044034	160.445548	23.428046	2.49832
min	1.000000	0.000000	3.000000	0.37000
25%	19.300000	0.000000	78.000000	4.26000
50%	43.500000	4.000000	93.000000	5.75500
75%	56.200000	28.000000	97.000000	7.49250
max	87.300000	2500.000000	99.000000	17.60000

	Diphtheria	HIV/AIDS	GDP	Population \
count	2919.000000	2938.000000	2490.000000	2.286000e+03
mean	82.324084	1.742103	7483.158469	1.275338e+07
std	23.716912	5.077785	14270.169342	6.101210e+07
min	2.000000	0.100000	1.681350	3.400000e+01
25%	78.000000	0.100000	463.935626	1.957932e+05
50%	93.000000	0.100000	1766.947595	1.386542e+06
75%	97.000000	0.800000	5910.806335	7.420359e+06
max	99.000000	50.600000	119172.741800	1.293859e+09

	thinness 1-19 years	thinness 5-9 years \
count	2904.000000	2904.000000
mean	4.839704	4.870317
std	4.420195	4.508882
min	0.100000	0.100000
25%	1.600000	1.500000
50%	3.300000	3.300000
75%	7.200000	7.200000
max	27.700000	28.600000

	Income composition of resources	Schooling
count	2771.000000	2775.000000
mean	0.627551	11.992793
std	0.210904	3.358920
min	0.000000	0.000000
25%	0.493000	10.100000
50%	0.677000	12.300000
75%	0.779000	14.300000
max	0.948000	20.700000

```
[16]: print("amount of missing data for each column:")
      print(np.sum(data.isnull()))
```

```
amount of missing data for each column:
Country      0
Year         0
Status       0
```

Life expectancy	10
Adult Mortality	10
infant deaths	0
Alcohol	194
percentage expenditure	0
Hepatitis B	553
Measles	0
BMI	34
under-five deaths	0
Polio	19
Total expenditure	226
Diphtheria	19
HIV/AIDS	0
GDP	448
Population	652
thinness 1-19 years	34
thinness 5-9 years	34
Income composition of resources	167
Schooling	163
dtype: int64	

```
[17]: data.fillna(data.mean(),inplace=True)
      print("No more missing data:")
      print(np.sum(data.isnull()))
```

```
No more missing data:
Country      0
Year         0
Status       0
Life expectancy  0
Adult Mortality  0
infant deaths  0
Alcohol      0
percentage expenditure  0
Hepatitis B   0
Measles      0
BMI          0
under-five deaths  0
Polio        0
Total expenditure  0
Diphtheria   0
HIV/AIDS     0
GDP          0
Population   0
thinness 1-19 years  0
thinness 5-9 years  0
Income composition of resources  0
Schooling    0
```

dtype: int64

```
[18]: data.drop(columns = ["Hepatitis B", "Population"], inplace=True)
```

```
[19]: data["Status"]
```

```
[19]: 0      Developing
      1      Developing
      2      Developing
      3      Developing
      4      Developing
      ...
      2933    Developing
      2934    Developing
      2935    Developing
      2936    Developing
      2937    Developing
      Name: Status, Length: 2938, dtype: object
```

```
[20]: labels, levels = pd.factorize(data["Status"])
      data["Status"] = labels
```

```
[21]: data
```

```
[21]:
```

	Country	Year	Status	Life expectancy	Adult Mortality	\
0	Afghanistan	2015	0	65.0	263.0	
1	Afghanistan	2014	0	59.9	271.0	
2	Afghanistan	2013	0	59.9	268.0	
3	Afghanistan	2012	0	59.5	272.0	
4	Afghanistan	2011	0	59.2	275.0	
...	...	...		...	...	
2933	Zimbabwe	2004	0	44.3	723.0	
2934	Zimbabwe	2003	0	44.5	715.0	
2935	Zimbabwe	2002	0	44.8	73.0	
2936	Zimbabwe	2001	0	45.3	686.0	
2937	Zimbabwe	2000	0	46.0	665.0	

	infant deaths	Alcohol	percentage expenditure	Measles	BMI	\
0	62	0.01	71.279624	1154	19.1	
1	64	0.01	73.523582	492	18.6	
2	66	0.01	73.219243	430	18.1	
3	69	0.01	78.184215	2787	17.6	
4	71	0.01	7.097109	3013	17.2	
...	...	...	...	...	...	
2933	27	4.36	0.000000	31	27.1	
2934	26	4.06	0.000000	998	26.7	
2935	25	4.43	0.000000	304	26.3	

2936	25	1.72	0.000000	529	25.9
2937	24	1.68	0.000000	1483	25.5

	under-five deaths	Polio	Total expenditure	Diphtheria	HIV/AIDS \
0	83	6.0	8.16	65.0	0.1
1	86	58.0	8.18	62.0	0.1
2	89	62.0	8.13	64.0	0.1
3	93	67.0	8.52	67.0	0.1
4	97	68.0	7.87	68.0	0.1
...	...	...	...	...	...
2933	42	67.0	7.13	65.0	33.6
2934	41	7.0	6.52	68.0	36.7
2935	40	73.0	6.53	71.0	39.8
2936	39	76.0	6.16	75.0	42.1
2937	39	78.0	7.10	78.0	43.5

	GDP	thinness 1-19 years	thinness 5-9 years \
0	584.259210	17.2	17.3
1	612.696514	17.5	17.5
2	631.744976	17.7	17.7
3	669.959000	17.9	18.0
4	63.537231	18.2	18.2
...	...	...	...
2933	454.366654	9.4	9.4
2934	453.351155	9.8	9.9
2935	57.348340	1.2	1.3
2936	548.587312	1.6	1.7
2937	547.358879	11.0	11.2

	Income composition of resources	Schooling
0	0.479	10.1
1	0.476	10.0
2	0.470	9.9
3	0.463	9.8
4	0.454	9.5
...	...	...
2933	0.407	9.2
2934	0.418	9.5
2935	0.427	10.0
2936	0.427	9.8
2937	0.434	9.8

[2938 rows x 20 columns]

```
[22]: #the name of the columns has to be found by using data.columns first
data.rename(columns={" thinness 1-19 years":
    ↳"thinness_10-19_years"},inplace=True)
```

[23]: data

```
[23]:      Country  Year  Status  Life expectancy  Adult Mortality  \
0      Afghanistan  2015      0          65.0          263.0
1      Afghanistan  2014      0          59.9          271.0
2      Afghanistan  2013      0          59.9          268.0
3      Afghanistan  2012      0          59.5          272.0
4      Afghanistan  2011      0          59.2          275.0
...      ...      ...      ...      ...      ...
2933      Zimbabwe  2004      0          44.3          723.0
2934      Zimbabwe  2003      0          44.5          715.0
2935      Zimbabwe  2002      0          44.8           73.0
2936      Zimbabwe  2001      0          45.3          686.0
2937      Zimbabwe  2000      0          46.0          665.0

      infant deaths  Alcohol  percentage expenditure  Measles  BMI  \
0              62      0.01          71.279624      1154  19.1
1              64      0.01          73.523582       492  18.6
2              66      0.01          73.219243       430  18.1
3              69      0.01          78.184215      2787  17.6
4              71      0.01           7.097109      3013  17.2
...      ...      ...      ...      ...      ...
2933              27      4.36          0.000000        31  27.1
2934              26      4.06          0.000000       998  26.7
2935              25      4.43          0.000000       304  26.3
2936              25      1.72          0.000000       529  25.9
2937              24      1.68          0.000000      1483  25.5

      under-five deaths  Polio  Total expenditure  Diphtheria  HIV/AIDS  \
0              83      6.0          8.16          65.0      0.1
1              86     58.0          8.18          62.0      0.1
2              89     62.0          8.13          64.0      0.1
3              93     67.0          8.52          67.0      0.1
4              97     68.0          7.87          68.0      0.1
...      ...      ...      ...      ...      ...
2933              42     67.0          7.13          65.0     33.6
2934              41      7.0          6.52          68.0     36.7
2935              40     73.0          6.53          71.0     39.8
2936              39     76.0          6.16          75.0     42.1
2937              39     78.0          7.10          78.0     43.5

      GDP  thinness_10-19_years  thinness 5-9 years  \
0    584.259210          17.2          17.3
1    612.696514          17.5          17.5
2    631.744976          17.7          17.7
3    669.959000          17.9          18.0
4     63.537231          18.2          18.2
```

...	...	...	...
2933	454.366654	9.4	9.4
2934	453.351155	9.8	9.9
2935	57.348340	1.2	1.3
2936	548.587312	1.6	1.7
2937	547.358879	11.0	11.2

	Income composition of resources	Schooling
0	0.479	10.1
1	0.476	10.0
2	0.470	9.9
3	0.463	9.8
4	0.454	9.5
...	...	...
2933	0.407	9.2
2934	0.418	9.5
2935	0.427	10.0
2936	0.427	9.8
2937	0.434	9.8

[2938 rows x 20 columns]

```
[24]: data.columns
```


```
[24]: Index(['Country', 'Year', 'Status', 'Life expectancy ', 'Adult Mortality',
        'infant deaths', 'Alcohol', 'percentage expenditure', 'Measles ',
        ' BMI ', 'under-five deaths ', 'Polio', 'Total expenditure',
        'Diphtheria ', ' HIV/AIDS', 'GDP', 'thinness_10-19_years',
        ' thinness 5-9 years', 'Income composition of resources', 'Schooling'],
        dtype='object')
```

```
[25]: grouped = data.groupby('Country')
fig = plt.figure(figsize=(30,10))
ax = plt.axes()
grouped.mean()['Life expectancy '].plot(ax = ax)
```

```
/home/rom/Desktop/AIT/Programming/lib/python3.6/site-
packages/pandas/plotting/_matplotlib/core.py:1235: UserWarning: FixedFormatter
should only be used together with FixedLocator
    ax.set_xticklabels(xticklabels)
```

```
[25]: <AxesSubplot:xlabel='Country'>
```






```
[26]: #print(grouped.mean()['Life expectancy '])
series = grouped.mean()['Life expectancy ']
series.sort_values()
```

```
[26]: Country
Sierra Leone          46.11250
Central African Republic 48.51250
Lesotho                48.78125
Angola                 49.01875
Malawi                 49.89375
...
France                82.21875
Switzerland           82.33125
Iceland               82.44375
Sweden               82.51875
Japan                 82.53750
Name: Life expectancy , Length: 193, dtype: float64
```

```
[27]: grouped = data.groupby("Status")
grouped.mean()['Life expectancy '].plot()
print(levels)
```

```
Index(['Developing', 'Developed'], dtype='object')
```



```
[28]: from scipy import stats
data.groupby("Status").apply(lambda df: stats.ttest_ind(df['Life expectancy '],
↳df['Status']))
```


```
[28]: Status
0    (367.7986278092848, 0.0)
1    (450.1250388656375, 0.0)
dtype: object
```

```
[29]: col = list(data.columns)
```

```
[30]: import seaborn as sns


sns.pairplot(data,x_vars=["Life expectancy "],y_vars = col)
#adult mortality, schooling, income composition seem to be the most correlated
```

```
[30]: <seaborn.axisgrid.PairGrid at 0x7f0f0a5cd7b8>
```



```
[31]: data['Life expectancy '].hist()  
      #its a skewed graph so it's not normal
```

```
[31]: <AxesSubplot:>
```



```
[32]: col
```

```
[32]: ['Country',  
      'Year',  
      'Status',  
      'Life expectancy ',  
      'Adult Mortality',  
      'infant deaths',  
      'Alcohol',  
      'percentage expenditure',  
      'Measles ',  
      ' BMI ',  
      'under-five deaths ',  
      'Polio',  
      'Total expenditure',  
      'Diphtheria ',  
      ' HIV/AIDS',  
      'GDP',
```

```
'thinness_10-19_years',
' thinness 5-9 years',
'Income composition of resources',
'Schooling']
```

```
[33]: #From the question above, these three columns seem to have the best linear
      ↪ correlation to life expectancy
X = data.loc[:,["Adult Mortality","Income composition of
      ↪resources","Schooling"]].values.astype(float)
X_all = data.drop(columns=['Country','Life expectancy ']).values.astype(float)
y = data['Life expectancy '].values.astype(float)

print(X.shape)
print(X_all.shape)

#Standardizing the data
mean = np.mean(X,axis=0)
std = np.std(X,axis=0)
X_norm = (X-mean)/std

mean_all = np.mean(X_all,axis=0)
std_all = np.std(X_all,axis=0)
X_all_norm = (X_all-mean_all)/std_all

#Shuffling the data
ix = np.arange(X.shape[0])
np.random.shuffle(ix)
m = X.shape[0]
percentage = 0.7
ix_train = ix[:int(m*percentage)]
ix_test = ix[int(m*percentage):]

X_norm_train = X_norm[ix_train]
X_all_norm_train = X_all_norm[ix_train]
X_norm_test = X_norm[ix_test]
X_all_norm_test = X_all_norm[ix_test]
y_train = y[ix_train]
y_test = y[ix_test]

print(X_norm_test.shape[0]/(X_norm_train.shape[0]+X_norm_test.shape[0]))

#Checking whether the shapes are in the correct form
assert X_norm_test.shape[0]/(X_norm_train.shape[0]+X_norm_test.shape[0]) < 0.31
      ↪and X_norm_test.shape[0]/(X_norm_train.shape[0]+X_norm_test.shape[0]) > 0.29
```

```
(2938, 3)
```

```
(2938, 18)
```

0.30020422055820284

```
[34]: class Regression():
    def
    ↪__init__(self,X,y,grad_method,max_iter,alpha,tol,decay,decay_iter,decay_rate,stop_delay_cou
    ↪

        self.X = X
        self.y = y
        self.grad_method = grad_method
        self.max_iter = max_iter
        self.alpha = alpha
        self.tol = tol
        self.decay = decay
        self.decay_iter = decay_iter
        self.decay_rate = decay_rate
        self.stop_delay_counter = stop_delay_counter
        self.verbose = verbose
        self.lam = lam
        self.poly = poly
        self.poly_deg = poly_deg

    #1. hypothesis function
    def h(self, theta):
        hypothesis = self.X@theta
        return hypothesis

    #2. cost function
    def cost(self, theta):
        J = 1/(2*X.shape[0])*(self.h(X,theta)-y).T@(self.h(X,theta)-y)
        return J

    def cost_reg(self, X, y, theta,lamb):
        J = 1/(2*X.shape[0])*(self.h(X,theta)-y).T@(self.h(X,theta)-y) +
    ↪lamb*np.sum(theta@theta)
        return J

    #3. gradient function
    def gradient_reg(self, X, y, theta, average = False):
        dJ = X.T@(self.h(X,theta)-y)/(X.shape[0]) + theta*lamb
        return dJ

    def gradient(self, X, y, theta, average = False):
        dJ = X.T@(self.h(X,theta)-y)/(X.shape[0])
        return dJ

    def mini_batch():
        cost = []
```

```

theta = initial_theta
iteration = 0
cost.append(self.cost(X,y,theta,average))
ix = np.arange(y.size)
percentage = 0.1
ix_range = int(percentage*ix.size)
for n in range(max_iteration):
    np.random.shuffle(ix)
    X_batch = X[ix_range]
    y_batch = X[ix_range]
    gradient = self.gradient(X_batch,y_batch,theta,average)
    theta = theta - alpha*gradient
    cost.append(self.cost(X_batch,y_batch,theta,average))
    iteration += 1
cost = np.array(cost)
return theta,cost,iteration

def stochastic():
    cost = []
    theta = initial_theta
    iteration = 0
    cost.append(self.cost(X,y,theta,average))
    for n in range(max_iteration):
        for i in X.shape[0]:
            if(self.grad_method == 'ridge'):
                gradient = self.gradient_reg(X[i],y,theta,average)
            else:
                gradient = self.gradient(X[i],y,theta,average)
            theta = theta - alpha*gradient
            cost.append(self.cost(X[i],y,theta,average))
            iteration += 1
    cost = np.array(cost)
    return theta,cost,iteration

def batch():
    cost = []
    theta = initial_theta
    iteration = 0
    cost.append(self.cost(X,y,theta,average))
    for n in range(max_iteration):
        gradient = self.gradient(X,y,theta,average)
        theta = theta - alpha*gradient
        cost.append(self.cost(X,y,theta,average))
        iteration += 1
    cost = np.array(cost)
    return theta,cost,iteration

```

```

def polynomial_features():
    pass

def decay_learning_rate():
    pass

def ridge():
    pass

def fit():
    pass

def predict():
    pass

def score():
    pass

def mse():
    pass

#would've been better if the class wasn't fixed to the form Regression(X, y,
→grad_method, max_iter, alpha, tol, decay, decay_iter, decay_rate,
→stop_delay_counter, verbose, lam, poly, poly_deg)
#because i have written a class but not in this way

```

2). **Regression** - Prepare your X and y into Numpy array (you have to map from Pandas to numpy). For X, prepare two versions of them. For first X\_selected, you have to choose the most 3 important features from above, and for second X\_all, simply use all features (you may want to omit Country since they are categorical). Set y to life expectancy. (1 or 0pt)

- Perform standardization using Numpy way (NOT sklearn way). (1 or 0pt)
- Perform train-test split by using Numpy way (NOT sklearn way). Use test size of 0.3. (1 or 0pt)
- Perform assertion whether your splitting is correct accordingly (1 or 0pt)
- Write a class Regression(X, y, grad\_method, max\_iter, alpha, tol, decay, decay\_iter, decay\_rate, stop\_delay\_counter, verbose, lam, poly, poly\_deg) that can perform the followings:
  - Mini-batch, Stochastic, and Batch Gradient Descent (each 2pts)
  - Polynomial of degree k (2 or 0pt)
  - Decay learning rate (1 or 0pt)
    - \* Decay learning rate is a learning rate that becomes smaller after certain iteration. For example, after 5 iterations, the learning rate will reduce to 95% of the current learning rate.
    - \* To implement it, simply multiply current learning rate with some constant decay\_rate. For now, set it to 0.9



- Regularization with ridge (2 or 0pt)
- Must have at least four methods for `fit()` (i.e., for finding weights) `predict()` (i.e., for predicting  $X_{\text{test}}$  data), `score()` (i.e., for returning  $r^2$  score), and `mse()` (return mse) (each 1pt)
- Accepts  $X$ ,  $y$ , `grad_method` (default set to "batch"), `alpha` (learning rate), `max_iter`, `tol`, `decay` (whether to use decay learning rate; default set to False), `decay_iter` (after how many iterations will the decay apply), `stop_delay_counter` (this is the maximum number of times that decay the learning rate), `verbose` (default is set to False, whether model will display the Cost for each iteration), `lam` (this is the ridge regularization parameter), `poly` (default is set to False), and `poly_deg` (default is set to 2) (each 1/13pt)
- Create the following 3 models **from your class** (For any unspecified parameters, feel free to use any :D)
  1. For the first model, transform your feature using polynomial degree 3, then perform linear regression with batch gradient descent with early stopping of tol 1e-3 (1 or 0pt)
  2. For the second model, perform linear regression with mini-batch gradient descent with early stopping of tol 1e-3 (1 or 0pt)
  3. For the third model, perform ridge regression with stochastic gradient descent with early stopping of tol 1e-3 and decay set to True and lam to 1e-4 (1 or 0pt)
- Create Lasso model from Sklearn with default parameters (1 or 0pt)
- For these four models, using two different versions of  $X$ , perform a cross validation of 10 folds, comparing the four models \* two versions of  $X$ . Here you should implement cross validation. Report which one is the best candidate model (3pts for implement from scratch or 1pt for using sklearn)
  - Recall that in a 10 folds cross validation, you split your data into 10 even pieces. Then you run 10 iterations, where in each iteration, you pick 1 of this piece as the validation set, and the rest as training set. Once you reach the 10th iteration, you would have already exhaust all the 10 pieces as validation set.
- Using the best model, fit again with the training data. Plot the weights using bar charts along the feature names. Before you actually plot the weights, we need to multiply these weights by their feature standard deviation, so to reduce these weights to same unit of measure. Interpret these weights and what they imply. (For those who are curious why we need to multiply with std, you may read this > [https://scikit-learn.org/stable/auto\\_examples/inspection/plot\\_linear\\_model\\_coefficient\\_interpretation.html#interpretation-coefficients-scale-matters](https://scikit-learn.org/stable/auto_examples/inspection/plot_linear_model_coefficient_interpretation.html#interpretation-coefficients-scale-matters)) (2 or 0pt)
- Perform predictions on testing data. Print adjusted  $r^2$  and mse. (1 or 0pt)
- Plot the predicted values against actual values (1 or 0pt)

### 3). Classification

- Change your  $y$  to discrete value. Here split  $y$  into three class,  $\{0, 1, 2\}$ , where 0 belongs to low life expectancy group, and 2 for the high life expectancy group. (1 or 0pt)
- Write a class for multinomial logistic regression with stochastic gradient descent. Must have at least six methods for `fit()` (i.e., for finding weights) `predict()` (i.e., for predicting  $X_{\text{test}}$  data), `accuracy()` (i.e., for returning accuracy score), `recall()`, `precision()`, and `f1()` (each 1pt)

- Using the best `X_train` of the two suggested by the cross validation step, fit the data with your class. (1 or 0pt)
- Perform predictions on testing data. Print accuracy, recall, precision, and `f1_score` from your class. (1 or 0pt)
- Plot the decision boundary with the `X_test` data. To plot this, you may want to choose only 2 features. (1 or 0pt)

4). **Final verdict**

- Attempt to do whatever ways - including sklearn or scratch - or change your features, or do feature engineering such that your mse is lowest possible. (0 to 5pts - following class normal distributions)